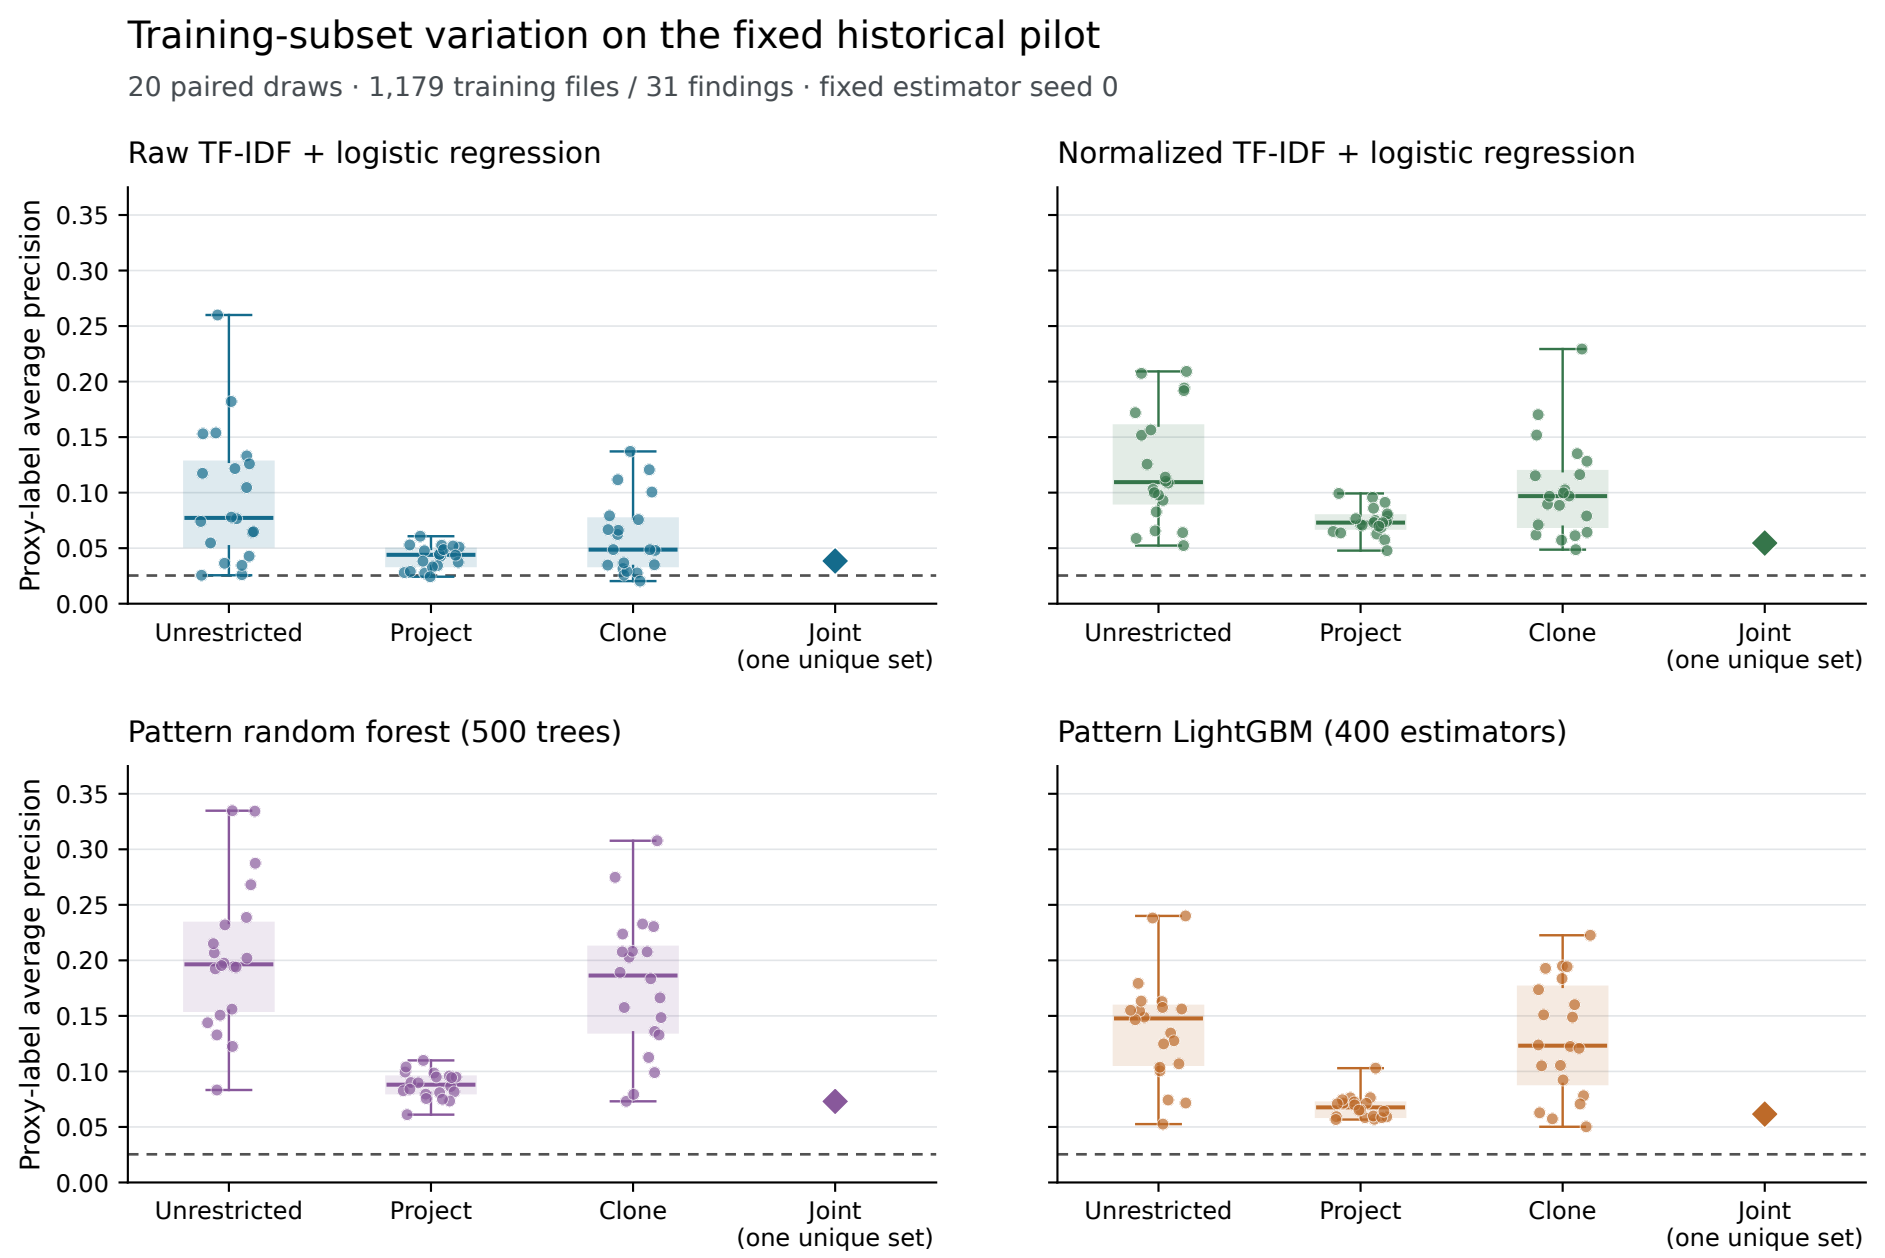
Training-subset variation on the fixed historical pilot
20 paired draws · 1,179 training files / 31 findings · fixed estimator seed 0
Raw TF-IDF + logistic regression
Normalized TF-IDF + logistic regression
Pattern random forest (500 trees)
Pattern LightGBM (400 estimators)
Proxy-label average precision
0.00
0.05
0.10
0.15
0.20
0.25
0.30
0.35
Unrestricted
Project
Clone
Joint
(one unique set)
Dashed line: exact random-ranking mean AP = 0.02536. Ranges are descriptive, not confidence intervals.